

4G System Handover Optimisation

Iain Cuthbertson

201015895

Computer & Electronic Systems

University of Strathclyde

April 2014

I hereby declare that this work has not been submitted for any other degree/course at this University or any other institution and that, except where reference is made to the work of other authors, the material presented is original and entirely the result of my own work at the University of Strathclyde under the supervision of Dr. Robert C. Atkinson.

Signature _____

Abstract

With more and more customers using mobile communications it is important for the service providers to give their customers the best Quality-of-Service (QoS) they can. Many providers have taken to improve their networks and make them more appealing to customers. One such improvement that providers can give to their customers is to improve the reliability of their network meaning that customers calls are less likely to be dropped by the network.

This dissertation explores improving the reliability of a 4G network by optimising the parameters used in handovers. The process of handover within mobile communication networks is very important and allowing for users to move around freely while still staying connected to the network. The parameters used in the handover process are the Time-to-Trigger (TTT) and Hysteresis (hys). These parameters are used to determine where a base station is better than the serving base station by enough to warrant a handover taking place. The challenge in optimising the handover parameters is that there is a fine balance that needs to be struck between calls being dropped due to a handover failing and the connection switching back and forth between two base stations, unnecessarily, wasting the network's resources. The approach taken is to use a machine learning technique known as Q-Learning to optimise the handover parameters by generating a policy that can be followed to adjust the parameters as needed.

Acknowledgements

Contents

Abstract	ii
Acknowledgements	iii
Nomenclature	vi
List of Figures	ix
List of Tables	x
List of Algorithms	xi
1 Introduction	1
1.1 Project Objectives	1
1.2 Project Outcomes	2
1.3 Dissertation Outline	2
2 LTE	3
2.1 Network Structure	3
2.2 Self-Organising Network	5
2.3 Handover	5
2.3.1 Parameters	6
2.3.2 Procedure	7
3 Machine Learning	10
3.1 Reinforcement Learning	11
3.2 Q-Learning	12
3.2.1 Example	14
4 Simulation Design	15
4.1 Simulation Parameters	16
4.2 Discrete Event Simulation	16
4.3 Mobility Model	16

4.3.1	Random Direction Model	16
4.3.2	Random Waypoint Model	17
4.4	Propagation Model	17
4.4.1	Okumura-Hata Model	18
4.4.2	Egli Model	18
4.4.3	Cost231-Hata Model	19
4.5	Simulation Testing	20
5	Handover Parameter Optimisation	21
5.1	Approach	21
5.1.1	Assessment Process	23
5.2	Walking Speed	24
5.2.1	Large Starting Values	24
5.2.2	Middle Starting Values	26
5.2.3	Small Starting Values	28
5.2.4	Large hys and Small TTT Starting Values	30
5.3	Vehicle Speed	31
5.3.1	Large Starting Values	32
5.3.2	Middle Starting Values	33
5.3.3	Small Starting Values	35
5.3.4	Large hys and Small TTT Starting Values	37
5.4	Evaluation	38
6	Future Work	40
7	Conclusions	41
	Bibliography	42
	Appendices	44
A	System Design and Evaluation	45

Nomenclature

3G	Third Generation
4G	Fourth Generation
AI	Artificial Intelligence
dB	Decibels
DES	Discrete Event Simulation
DL	Downlink
eNodeB	Evolved Node B
EPC	Evolved Packet Core
E-UTRAN	Evolved UMTS Terrestrial Radio Access Network
HSS	Home Subscriber Server
hys	Hysteresis
LTE	Long Term Evolution
MME	Mobility Management Entity
P-GW	Packet Data Network Gateway
PCell	Primary Cell
QoS	Quality-of-Service
RSS	Received Signal Strength
SCell	Secondary Cell
S-GW	Serving Gateway
TTT	Time-to-Trigger

UE	User Equipment
UL	Uplink
UMTS	Universal Mobile Telecommunications System

List of Figures

2.1	Illustration of E-UTRAN.	4
2.2	Illustration of EPC.	5
2.3	Graph of TTT values.	7
4.1	Graph of different Propagation Models.	20
5.1	Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=5.12s hys=10dB when UE traveling at walking speeds. . . .	25
5.2	Illustration of how the TTT and hys values changed over time for large values when UE traveling at walking speeds.	26
5.3	Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0.256s hys=5dB when UE traveling at walking speeds. . . .	27
5.4	Illustration of how the TTT and hys values changed over time for medium values when UE traveling at walking speeds.	28
5.5	Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0s hys=0dB when UE traveling at walking speeds.	29
5.6	Illustration of how the TTT and hys values changed over time for medium values when UE traveling at walking speeds.	29
5.7	Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0.08s hys=7.5dB when UE traveling at walking speeds. . .	30
5.8	Illustration of how the TTT and hys values changed over time for medium values when UE traveling at vehicle speeds.	31
5.9	Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=5.12s hys=10dB when UE traveling at walking speeds. . . .	32
5.10	Illustration of how the TTT and hys values changed over time for large values when UE traveling at vehicle speeds.	33
5.11	Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0.256s hys=5dB when UE traveling at vehicle speeds. . . .	34
5.12	Illustration of how the TTT and hys values changed over time for medium values when UE traveling at vehicle speeds.	35
5.13	Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0s hys=0dB when UE traveling at vehicle speeds.	36

5.14	Illustration of how the TTT and hys values changed over time for medium values when UE traveling at vehicle speeds.	36
5.15	Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0.08s hys=7.5dB when UE traveling at vehicle speeds. . . .	37
5.16	Illustration of how the TTT and hys values changed over time for medium values when UE traveling at vehicle speeds.	38

List of Tables

2.1	Table of the different LTE hys values.	6
2.2	Table of the different LTE TTT values.	7
2.3	Table of the different LTE Trigger types and their criteria.	8

List of Algorithms

3.1	Q-Learning Procedure.	13
-----	-------------------------------	----

Chapter 1

Introduction

Handover is a key process in Fourth Generation (4G) cellular networks. It allows for customers to be able to continue their calls and using services while moving out of range of the base station they are currently connected to. A reinforcement learning technique known as Q-Learning was used to carry out the optimisation of the handover parameters.

The performance of handovers within a cellular network is very important. Initiating a handover too late can cause an on-going call to be dropped by the network. This would be viewed as poor Quality-of-Service (QoS) by the customer. However, initiating a handover too early also comes with its own problems. It can cause handover ping-pong's where a mobile's connection switches back and forth between base stations in a short period of time and this wastes resources within the network.

It was also required that a simulation was created that would simulate mobile phones moving around a group of network base stations. This simulation would be used to determine the performance of the machine learning algorithm compared to if no optimisation was taking place. It would also supply the machine learning algorithm with rewards to tell it if it is making decisions that are improving the performance or making it worse.

1.1 Project Objectives

The main objective of this project was to develop a machine learning algorithm to fine tune the parameters used in the handover process of 4G networks.

Key Objectives:

- Research and understand the parameters used in 4G handovers.
- Create a basic simulation of a 4G network with mobiles moving around a group of base stations.

- Implement a machine learning algorithm to fine tune the handover parameters to improve the performance of the network.
- Evaluate the success of the machine learning approach in 4G handovers and compare this to using static parameters

1.2 Project Outcomes

During the time spent on the project all of the objectives have been met. The parameters that are used in the handover process, the Time-to-Trigger (TTT) and Hysteresis (hys), were researched and the ways they are used were discovered. A basic simulation of a cellular network was created and a Q-Learning algorithm was developed that optimised the handover parameters.

Experiments were run to find the worth of the machine learning system against using static values for the handover parameters. It was found that the machine learning system was able to improve the performance of the network. However, the system was found to be capable of having the potential of performing even better.

1.3 Dissertation Outline

The structure of this dissertation is as follows: Chapter 2 describes a basic account of a 4G network along with the handover process. Chapter 3 gives a brief discussion on the different types of machine learning and goes more in-depth into reinforcement learning and Q-Learning. Chapter 4 describes the different propagation and mobility models considered for use within the simulation, as well as giving accounts of how the simulation was designed and the testing done to make sure the simulation was functioning as required. Chapter 5 describes the approach taken to the machine learning problem and the experiments run to find the worth of the system. This chapter also gives the results of the experiments and these results are discussed. Chapter 6 discusses the possible future work that could take this project further. Finally, Chapter 7 gives an account of the conclusions drawn from the project.

Chapter 2

LTE

Mobile communications is on to its Fourth Generation (4G) of network infrastructure with Long Term Evolution (LTE). This network infrastructure was developed by 3GPP and is an improvement upon Universal Mobile Telecommunications System (UMTS), which is a Third Generation (3G) network. LTE has Downlink (DL) speeds of up to 300Mbit/s and Uplink (UL) speeds of 75Mbit/s with a varying bandwidth that ranges from 1.4MHz to 20MHz . The development of faster downloads speed was driven by the consumers want for better quality images, faster Internet browsing and smoother video streaming. [1]

2.1 Network Structure

The structure of the LTE network can be broken down into 3 main parts, the User Equipment (UE), the Evolved UMTS Terrestrial Radio Access Network (E-UTRAN) and the Evolved Packet Core (EPC). The UE can simply be considered as a standard mobile phone or smartphone. The purpose of the E-UTRAN is to connect a UE to the EPC and is made up of just one component, the Evolved Node B (eNodeB) or simply put base station. Figure 2.1 shows an illustration of the E-UTRAN. An UE will only communicate with one eNodeB at any time and this eNodeB is known as the Serving eNodeB. An eNodeB proves two main functions within the network; the first is to send all the radio traffic for an UE on the DL as well as receiving any traffic sent from the UE on the UL. The second function of the eNodeB is to control low-level operations such as handovers as well as provide the signalling for such operations. Due to the eNodeB's having the added complexity of controlling operations such as handovers it moves more of the processing from the core network to the edge of the network, reducing the latency of decisions being made. Every eNodeB is connected to the EPC using the S1 interface. They can also be connected to other eNodeB's by the X2 interface. This interface is mainly used for signalling and forwarding data from a

serving eNodeB to a neighbouring eNodeB during a handover.




Figure 2.1: Illustration of E-UTRAN.

An illustration of an EPC can be seen in Figure 2.2. It can be seen that the EPC is made up of four main components, they are the: Mobility Management Entity (MME), Home Subscriber Server (HSS), Packet Data Network Gateway (P-GW) and Serving Gateway (S-GW). The MME provide the high-level operations for a UE such as security and managing non-radio communication data streams as well as controlling other elements within the EPC. There are very few MME's within the LTE network, generally assigned to a certain geographical region. A UE will be assigned to a single MME known as the Serving MME. The HSS is a central database that holds all the information about the network subscribers such as authentication and billing information. The P-GW is the bridge between the EPC and other packet data networks such as the Internet. The P-GW uses the SGi interface to exchange data with outside networks, such as the network operator's servers. The S-GW function is to forward data to and from the eNodeB's to P-GW, this means that the S-GW effectively acts like a router. Much like the MME a UE will be assigned to one S-GW and there will be very few within the network as a whole. [1, 2]




Figure 2.2: Illustration of EPC.

2.2 Self-Organising Network

LTE, just like any other mobile network, needs to be managed. Due to LTE being more complex than its predecessors the management of the network has also become more complex. The use of automation can simplify management of a network greatly. A technology that can be employed in LTE is that of a Self-Organising Network. This system provides three main functions, Self-configuration, Self-optimisation and Self-healing. [3, 4]

- **Self-configuration** allows for newly installed eNodeB to be configured automatically with the basic parameters for operation.
- **Self-optimisation** is a process where measurements from the UE and eNodeB along with performance measurements are used to optimise the network to make it perform better.
- **Self-healing** is designed to detect and identify failures within the network. After the self-healing system detects a failure it aims to recover from the failure without the need for human interaction with the system.

2.3 Handover

The process of handover is very important in mobile telecommunications. It involves moving the resource allocation for a mobile phone or a piece of UE from

one base station to another. This process is used to provide more Quality-of-Service (QoS) to customers by allowing them to continue to use provided services even after moving out of range of the original serving base station. To keep with the QoS it is important that handovers are done fast, have little-to-no disruption to the users experience and are completed with a very high success rate. If a handover is unsuccessful it is likely that an on going call will be dropped due to there not being enough resources available on a base station, known as an eNodeB in LTE, or the received signal strength to the UE drops below a certain threshold needed to maintain the call. Handovers are stated to take roughly 0.25 seconds to complete after the decision has been made for a handover to take place [5].

2.3.1 Parameters

In LTE there are two main parameters that are used in the handover process. These parameters are the Time-to-Trigger (TTT) and Hysteresis (hys). The hys is used to define how much better the Received Signal Strength (RSS) of a neighbouring base station must be than the serving base station for a handover to be considered. The values of hys are defined in Decibels (dB) and range from 0 to 10 dB in 0.5 dB increments, this results in there being 21 different values of hys. The full range of hys values can be seen in Table 2.1.

Index	0	1	2	3	4	5	6	7	8	9	10
hys (dB)	0.0	0.5	1.0	1.5	2.0	2.5	3.0	3.5	4.0	4.5	5.0
Index	11	12	13	14	15	16	17	18	19	20	
hys (dB)	5.5	6.0	6.5	7.0	7.5	8.0	8.5	9.0	9.5	10	

Table 2.1: Table of the different LTE hys values.

The TTT is a length of time, defined in seconds, that is used to define how long a neighbouring base station must look better than the serving base station for. There are 16 different values of TTT ranging from 0 to 5.12 seconds. Unlike with hys the TTT values do not increment linearly, instead they increment exponential with smaller increases at the lower values and bigger increases at the larger values. The full list of TTT values can be seen in Table 2.2 and a graph of how the TTT values increase can be seen in Figure 2.3.

Index	0	1	2	3	4	5	6	7	8
TTT (s)	0.0	0.04	0.064	0.08	0.1	0.128	0.16	0.256	0.32
Index	9	10	11	12	13	14	15		
TTT (s)	0.48	0.512	0.64	1.024	1.280	2.56	5.12		

Table 2.2: Table of the different LTE TTT values.




Figure 2.3: Graph of TTT values.

There are 336 different combinations of TTT and hys values. Having such a large range of combinations means that pairs of values can mean that a neighbouring eNodeB has to be better by a large hys but for a small TTT or vice-versa. This makes for an interesting dynamic for which pairs of values will work the best in any given environment.

2.3.2 Procedure

In LTE there are eight different triggers defined for initiating handovers within LTE. Table 2.3 shows different trigger events and how they are defined [6].

Event Type	Trigger Criteria
A1	Serving becomes better than a threshold.
A2	Serving becomes worse than a threshold.
A3	Neighbour becomes offset better than Primary Cell (PCell).
A4	Neighbour becomes better than threshold.
A5	PCell becomes worse than threshold1 and neighbour becomes better than threshold2.
A6	Neighbour becomes offset better than Secondary Cell (SCell).
B1	Inter RAT neighbour becomes better than threshold.
B2	PCell becomes worse than threshold1 and inter RAT neighbour becomes better than threshold2.

Table 2.3: Table of the different LTE Trigger types and their criteria.

Out of the eight triggers the A3 event is the most common and it is defined that a neighbouring eNodeB must give the UE better Received Signal Strength (RSS) by an amount defined by the hys, for a length of time defined by the TTT. An illustration of how the signals from the serving and neighbouring eNodeBs look in an A3 event can be seen in Figure (place holder). The A3 event can also be represented by the following equation:

$$RSS_{Neighbouring} + hys > RSS_{Serving} \quad (2.1)$$

Place holder for image of event A3

When a handover event is triggered a measurement report is sent from the UE to the serving eNodeB. The measurement report contains the information required for the serving eNodeB to make a decision on whether to initiate the handover or not. If the handover is triggered the serving eNodeB will signal to the neighbouring eNodeB that is to become the new serving eNodeB. If the serving eNodeB decides not to initiate the handover then no more action is taken until another handover event is triggered.

The full, high-level, procedure for a LTE handover is as follows:

1. If a neighbouring eNodeB is found to be better than the serving eNodeB a measurement report is sent by the UE to the serving eNodeB.
2. The serving eNodeB considers the information in the measurement report and decides whether or not a handover should take place.
3. If it is decided that a handover should take place then a message is sent to the neighbouring eNodeB to prepare resources for the UE.

4. Once the resources are ready for the UE the new serving eNodeB sends a message to the old eNodeB to release the resources it previously had for the UE
5. Finally a message is sent to the MME to finalise the handover process.

Chapter 3

Machine Learning

Machine learning is a form of Artificial Intelligence (AI) that involves designing and studying systems and algorithms with the ability to learn from data. This field of AI has many applications within research (such as system optimisation), products (such as image recognition) and advertising (such as adverts that use a users browsing history). There are many different paradigms that machine learning algorithms use. Algorithms can use training sets to train an algorithm to give appropriate outputs; other algorithms look for patterns in data; while others use the notion of rewards to find out if an action could be considered correct or not. [7] Three of the most popular types of machine learning algorithms are:

- **Supervised learning** is where an algorithm is trained using a training set of data. This set of data includes inputs and the known outputs for those inputs. The training set is used to fine-tune the parameters in the algorithm. The purpose of this kind of algorithm is to learn a general mapping between inputs and outputs so that the algorithm can give an accurate result for an input with an unknown output. This type of algorithm is generally used in classification systems.
- **Unsupervised learning** algorithms only know about the inputs they are given. The goal of such an algorithm is to try and find patterns or structure within the input data. Such algorithm would be given inputs and any patterns that are contained would become more and more common the more inputs the algorithm is given.
- **Reinforcement learning** uses an intelligent agent to perform actions within an environment. Any such action will yield a reward to the agent and the agent's goal is to learn about how the environment reacts to any given action. The agent then uses this knowledge to try and maximise its reward gains.

3.1 Reinforcement Learning

In reinforcement learning an intelligent agent is learning what action to do at any given time to maximise the notion of a reward. In the beginning the agent has no knowledge of what action it should take from any state within the learning environment. It must instead learn through trial and error, exploring all possible actions and finding the ones that perform the best.

The trade-off between exploration and exploitation is one of the main features of reinforcement and can greatly affect the performance of a chosen algorithm. A reinforcement learning algorithm must contemplate this trade-off of whether to exploit an action that resulted in a large reward or to explore other actions with the possibility of receiving a greater reward.

Another main feature of reinforcement learning is that the problem in question is taken into context as a whole. This is different from other types of machine learning algorithms, as they will not considered how the results of any sub-problems may affect the problem as a whole.

The basic elements required for reinforcement learning is as follows:

- A Model (M) of the environment that consists of a set of States (S) and Actions (A).
- A reward function (R).
- A value function (V).
- A policy (P).

The model of the environment is used to mimic the behaviour of the environment, such as predicting the next state and reward from a state and taken action. Models are generally used for planning by deciding what action to take while considering future rewards.

The reward function defines how good or bad an action is from a state. It is also used to define the immediate reward the agent can expect to receive. Generally a mapping between a state-action pair and a numerical value is used to define the reward that the agent would gain. The reward values are used to define the policy where the best value of state-action pair is used to define the action to take from a state.

While the reward function defines the immediate reward that can be gained from a state, the value function defines how good a state will be long-term. This difference can create possible conflicts of interest for an agent; so while its goal is to collect as much reward as possible, it has to weigh up the options of picking a state that may provide a lot of up front reward but not a lot of future reward

against a state with a lot of future reward but not a lot of immediate reward. This trade-off is similar to that of exploration versus exploit as it can define how successful a reinforcement algorithm is.

The policy is a mapping between a state and the best action to be taken from that state at any given time. Policies can be simple or complex; with a simple policy consisting of a lookup table, while more complex policies can involve search processes. In general most policies begin stochastic so that the agent can start to learn what actions are more optimal. [8]

3.2 Q-Learning

Q-Learning is a type of reinforcement learning algorithm where an agent tries to discover an optimal policy from its history of interactions from within an environment. What makes Q-Learning so powerful is that it will always learn the optimal policy for a problem regardless of the policy (which action a to take from a state s) it follows while learning as long as there is no limit on the number of times the agent can try an action. Due to this ability to always learn the optimal policy, Q-Learning is known as an Off-Policy learner. The history of interactions of an agent can be shown as a sequence of State-Action-Rewards:

$$< s_0, a_0, r_1, s_1, a_1, r_2, s_2, a_2 \dots >$$

This can be described as the agent was in State 0, did Action 0, received Reward 0 and transitioned into State 1; then did Action 1, received Reward 1 and transitioned into State 2; and so on.

The history of interactions can be treated as a sequence of experiences, with each experience being a tuple.

$$< s, a, r, s' >$$

The meaning of the tuple is that the agent was in State s , did Action a , received Reward r and transitioned in State s' . The experiences are what the agent uses to determine what the optimal action to take is at a given time.

The basic process of a Q-Learning algorithm can be seen in Algorithm 3.1. The general process requires that the learning agent is given a set of states, a set of actions, a discount factor γ and step size α . The agent also keeps a table of Q-Values, denoted by $Q(s, a)$ where s is a state and a is an action from that state. A Q-Value is also an average of all the experiences the agent has with a specific state-action pair. This allows for good and bad experiences to be averaged out to give a reasonable estimation of the actual value of state-action pair. The Q-Values are defined by Equation 3.1 where α is the step size which specifies how

much the new Q-Value is averaged with the old one, γ is the discount factor which specifies how much the agent considers the possible future rewards it will gain and the possible future rewards ($\max_{a'} Q(s', a')$) is the maximum of the Q-Values of all possible state-actions pairs from the action selected.

$$Q[s, a] = Q[s, a] + \alpha(r + \gamma \max_{a'} Q[s', a'] - Q[s, a]) \quad (3.1)$$

The table of Q-Values can either be initialised as empty or with some values pre-set to try and lead the agent to a specific goal state. Once the agent has initialised these parameters it observes the starting state. The starting state can either be chosen by random or be a pre-determined start state for the problem. The agent will then choose an action. Actions are chosen either stochastically or by a policy. Once an action has been chosen the agent will carry out the action and receive a reward. This reward is used to update the table of Q-Values using Equation 3.1. Finally the agent moves into the new state and repeats until termination; which can be either when the agent discovers a goal state or after a certain number of actions have been taken.

Require:

- S is a set of states
- A is a set of actions
- γ the discount reward factor
- α is the learning rate

```

1: procedure Q-LEARNING( $S, A, \gamma, \alpha$ )
2:   real array  $Q[S, A]$ 
3:   previous state  $s$ 
4:   previous action  $a$ 
5:   initialise  $Q[S, A]$  arbitrarily
6:   observe current state  $s$ 
7:   repeat
8:     select and carry out an action  $a$ 
9:     observe reward  $r$  and state  $s'$ 
10:     $Q[s, a] \leftarrow Q[s, a] + \alpha(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$ 
11:     $s \leftarrow s'$ 
12:  until termination
13: end procedure

```

Algorithm 3.1: Q-Learning Procedure.

After a Q-Learning algorithm has finished exploring the model of the environment it creates a policy. The policy is generated by searching across all actions for a state and finding the next state with the greatest value. The policy is therefore a lookup table that maps a state with the best possible next state. The policy created can then be used to solve the problem that the Q-Learning agent was exploring. [9]

3.2.1 Example

Chapter 4

Simulation Design

The simulation is a very important part of the project. It is required to provide the basic functionality of a LTE network. For simplicity the simulation was broken down into two main components; the mobile (UE) and the base station (eNodeB). Due to the project revolving around the handover process in LTE it made sense for the two main components on the simulation to be the mobile and base station as it is the mobile the triggers the measurement report and the base station that makes the decision on whether a handover should take place or not. Since the A3 event trigger is the most common I decided that it would be the only trigger implemented in the simulation to reduce the complexity within the simulation.

It was required for the mobile to be able to move freely around a group of base stations. It was decided that this movement should be random because if any machine learning algorithm can handle random movement then it should also be able to handle regimented movement. The movement that the mobile follows is defined by a Mobility Model and the choice of mobility model is explained in Section 4.3.

In wireless communications the received signal strength from a transmitter degrades the further away from the transmitter the receiver is. A propagation model can be used to define the way in which the signal strength degrades. The propagation model is very important to the simulation as it will defined how far away from a base station the mobile can be without dropping the call. A comparison and explanation of the choice of propagation model can be seen in Section 4.4.

4.1 Simulation Parameters

4.2 Discrete Event Simulation

The concept of time is very important in real world simulations. A popular method for creating this concept of time is Discrete Event Simulation (DES). It works on the basis of a scheduler where events, to be processed at a certain time, are passed to the scheduler and the scheduler then passes the event on to be processed. The way that DES creates the concept of time is by allowing events passed in the scheduler to have a time that it should be processed at, the scheduler will keep all the events that are currently still to be processed in an ordered list, the scheduler will then "jump" to the time the first event in the list is to be processed at and passes the event on to be processed. The process of "jumping" to the next time that the first event in the list is to be processed at is what gives DES the concept of time.

Another advantage to Discrete Event Simulation is that the passing of events from the scheduler not only gives a concept of time but it also allows for a message to be sent with the event to tell a different part of the simulation what to do at a given time. This message can also sent parameters for the other part of the simulation to use as well.

4.3 Mobility Model

A mobility model defines the way in which an entity will move. For the purposes of the simulation the mobility model used needed to random in nature. After some reason it was decided that mobility model to be used in the simulation would either be the Random Direction or Random Waypoint model.

4.3.1 Random Direction Model

The Random Direction Model is defined as follows:

1. Select a direction randomly between 0 and 355 degrees.
2. Select a random speed to move at.
3. Select a random duration to move for.
4. Move in the selected direction at the selected speed for the selected duration.
5. Repeat until termination.

An illustration of the movement given by the Random Direction Model can be seen in Figure (placeholder).

place holder for image of random direction

4.3.2 Random Waypoint Model

The Random Waypoint Model is defined as follows:

1. Randomly select the co-ordinates for a point within the environment.
2. Select a random speed to move at.
3. Select a random length of time to pause for when the destination is reached.
4. Move towards the selected co-ordinates at the selected speed
5. Pause for the randomly selected length of time.
6. Repeat until termination.

An illustration of the movement given by the Random Waypoint Model can be seen in Figure (placeholder).

place holder for image of random waypoint

It was decided that the Random Direction Model would be used in the simulation because the Random Waypoint Model has the problem that it is possible to select the co-ordinates of a point very close to where you begin and then pause for a long period time. The possibility of that happening isn't desired within the simulation. Random Direction does not have this problem and it is also possible to set boundaries on the parameters to make sure that a minimum distance is travelled. [10]

4.4 Propagation Model

A propagation model defines how the received signal from a transmitter decays the further from the transmitter you are. There are many different models available, all with different functions and purposes. After some research three models were considered; the Okumura-Hata Model, the Egli Model and the Cost231-Hata Model.

4.4.1 Okumura-Hata Model

The Okumura-Hata model is very popular for simulating transmissions in built up areas. Equations 4.1, 4.2 and 4.3 show the formulas for the model.

$$L_u = 69.55 + 26.16\log_{10}f - 13.82\log_{10}h_B - C_h + [44.9 - 6.55\log_{10}h_B]\log_{10}d \quad (4.1)$$

For small or medium sized cities:

$$C_H = 0.8 + (1.1\log_{10}f - 0.7)h_m - 1.56\log_{10}f \quad (4.2)$$

For large cities.

$$C_H = \begin{cases} 8.29(\log_{10}(1.54h_M))^2 - 1.1 & \text{if } 150 \leq f \leq 200 \\ 3.2(\log_{10}(11.75h_M))^2 - 4.97 & \text{if } 200 < f \leq 1500 \end{cases} \quad (4.3)$$

Where:

- L_u is the path loss (dB).
- H_B is the height of the base station antenna 30 to 200m.
- H_R is the height of the mobile antenna 1 to 10m.
- f is the frequency of the transmission 150 to 1500MHz.
- C_H is the antenna correction factor.
- d is the distance between the base station and the mobile 1 to 20km.

4.4.2 Egli Model

The Egli Model was another model that was considered for the simulation. Equations 4.4 and 4.5 shows the formulas for the model.

If h_m is less than or equal to 10m height:

$$P_L(dB) = 20\log_{10}f_c + 40\log_{10}d - 20\log_{10}h_b + 76.3 - 10\log_{10}h_m \quad (4.4)$$

If h_m is greater than or equal to 10m height:

$$P_L(dB) = 20\log_{10}f_c + 40\log_{10}d - 20\log_{10}h_b + 83.9 - 10\log_{10}h_m \quad (4.5)$$

Where:

- $P_L(dB)$ is the path loss (dB).

- h_B is the height of the base station antenna (m).
- h_M is the height of the mobile antenna (m).
- d is the distance between the base station and the mobile (m).
- f is the frequency of the transmission $3MHz$ to $3000MHz$.

4.4.3 Cost231-Hata Model

The Cost231-Hata model is an extension of the Okumura-Hata to work for frequencies between 1.5 GHz and 2 GHz. The formulas for this model can be seen in Equations 4.6, 4.7 and 4.8.

$$L = 46.3 + 33.9\log f - 13.82\log h_B - a(h_R) + [44.9 - 6.55\log h_B]\log d + C \quad (4.6)$$

$$a(h_R) = (1.1\log f - 0.7)h_M - (1.58\log f - 0.8) \quad (4.7)$$

$$C = \begin{cases} 0dB & \text{for medium cities and suburban areas} \\ 3dB & \text{for metropolitan areas} \end{cases} \quad (4.8)$$

Where:

- L is the path loss (dB).
- f is the frequency of the transmission 1500 to $2000MHz$.
- h_B is the height of the base station antenna 30 to $200m$.
- h_M is the height of the mobile antenna 1 to $10m$.
- d is the distance between the base station and the mobile 1 to $20km$.
- $a(h_R)$ is the antenna correction factor.

A comparison of the three propagation models can be seen in Figure 4.1. For the graphs the height of the mobile was said to be $1m$, while the height of the base station was said to be $60m$. The frequencies used in the models were $2000MHz$ for the Egli and Cost231-Hata models and $1500MHz$ for the Okumura-Hata model. These values for the propagation models are then taken away from a base station transmit power of $48dB$ so that the graphs are an estimate of the signal strength that would be received by a mobile at any distance up to $20km$.




Figure 4.1: Graph of different Propagation Models.

After comparing the three models above it was decided that the Cost231-Hata Model would be the one used in the simulation due to it working with frequencies up to 2000MHz (which is the minimum operating frequency of LTE) unlike the Okumura-Hata model which only works up to 1500MHz . The Cost231-Hata Model was also picked over the Egli Model because the Egli Model used more parameters that would add more complexity to the simulation. [11, 12]

4.5 Simulation Testing

Chapter 5

Handover Parameter Optimisation

5.1 Approach

The approach taken for optimising the handover parameters in LTE uses a Q-Learning algorithm based on the process given in Section 3.2. In the approach the model of the environment has a state for every combination of TTT and hys; giving a total number of 336 states. An action within the model can move to any other state that is different by one of the following changes to the handover parameters:

1. A single increase of TTT.
2. A single increase of hys.
3. A single increase of both TTT and hys.
4. A single decrease of TTT.
5. A single decrease of hys.
6. A single decrease of both TTT and hys.
7. A single value increase of TTT and a single value decrease of hys.
8. A single value increase of hys and a single value decrease of TTT.

For example if the learning agent is in the state where the TTT equals $0.256s$ and the hys equals $5.0dB$ and did action 3 from the list seen above; then the new TTT would equal $0.32s$ and the hys would equal $5.5dB$. The full list of hys values can be seen in Table 2.1 and the full list of TTT values can be seen in Table 2.2.

Having the actions only change the parameters by one increase or decrease of the TTT and hys values each time not only allows for more refined optimisation of the parameters but it also makes sure that no large changes can suddenly happen.

Due to the nature to the kind of problem that is being solved, the reward gained by an action is dynamic and is likely to be different each time it is taken. Rewards are based on the number of drop and ping-pong's accumulated in the simulation for current state in the environment model. The rewards are defined by the following equation:

$$Reward = Handover_{successful}(10 * Drops + 2 * PingPong's) \quad (5.1)$$

The coefficients in Equation 5.1 are given the values of 10 for drops and 2 for ping-pong's. Drops are extremely bad for the QoS of a communication system so it's given a large value and the reason ping-pong's are multiplied by 2 to remove the successful handover that was caused by the ping-pong and give the agent a penalty. The reward is given to the agent and the Q-Value for that state is updated just before agent selects the next action to take. The agent then selects new actions in discrete time steps, this allows for the simulation to run for fixed periods of time with TTT-hys pairs specified by a state in the environment model.

After the agent has been given enough time to try every action at least once the Q-Learning agent generates a policy. This policy can then be used to attempt to optimise the handover parameters by changing the TTT and hys values after a call is dropped or the connection ping-pongs between base stations. The Q-Learning agent still receives rewards every time a call is dropped or the connection ping-pong's while following the generated policy. Doing this allows for the system to always be learning; even after the initial learning process that generated the policy.

The optimisation system was tested in two scenarios. One scenario was to have 10 UE's moving randomly around 9 base stations, with each UE being 1 m in height and each base station being 60 m in height, with the layout as seen in Figure (Place Holder), using the Random Direction mobility model seen in Section 4.3, where the speed of the UE is 1 to 4m/s, which is walking speed and the duration of the direction is between 100 and 200 seconds. The other scenario is to have the UE moving at 10 to 15m/s, which is roughly 30mph. In these scenarios each mobile begins on top of one of the base stations before it starts moving so that handovers are not required as soon as the simulation starts. These scenarios would also be run with no fading in the RSS calculations so in to make the environment easier to learn for the agents.

Place holder for base station layout.

Each base station has its own Q-Learning agent to optimise the TTT and hys values for that specific base station. The agents are given 1000000 seconds to attempt to learn the environment that are working within, with each state being given 180 seconds to gain their reward. This length of time was chosen because there are 336 state each with a maximum of 8 actions, therefore the time needed to do all actions would take approximately 483840 seconds if each action was given 180 seconds. This length of time is less than half of the total time given so even due to the randomness of selecting next actions when learning the environment there should be enough time to try all state and most of the actions available. After the agents have learned the environment they generate a policy for their base station to follow. The simulation is then run for 200000 seconds to test how well the policies perform. The results for the scenarios can be seen in Section 5.1.1.

5.1.1 Assessment Process

The results are assessed by how well they performed compared to not making any changes to the TTT and hys values. The comparison is made by observing how to ratio of dropped calls and ping-pong's to successful handovers changes over time. The ratio of dropped calls is given by the following equation:

$$DropRatio = \frac{\#DroppedCalls}{\#SuccessfulHandovers} \quad (5.2)$$

The ratio for the connection ping-ponging between base stations is given by the following equation:

$$PingPongRatio = \frac{\#PingPong's}{\#SuccessfulHandovers + \#AttemptedHandovers} \quad (5.3)$$

It is also important to see how the TTT and hys values are being changes when the system is attempting to optimise them. This will allows it to be seen if the agents for the base stations come to some kind of consensus on the most optimal values of TTT and hys or if them come up with there own unique solutions. It is likely that Base Station 4 whose coverage overlaps with the coverage from every other Base Station, as seen in Figure (place holder), will come up with a different solution to the other Base Stations because it will be involved in the most handover attempts.

The simulation is run for four different starting states when compiling results for the UE moving at walking or vehicle speeds. The first starting state was to have the TTT and hys both start at their maximum values, 5.12 second and 10

dB respectively, to see how the system attempted to optimise the values as it is expected that a lot of dropped calls would occur for this set of values. The second start state was to give the TTT and hys their middle values, which are 0.256 seconds for TTT and 5 dB for hys. This start state would be expected to perform relatively well without any optimisation taking place due to the values neither being very large or very small. The third start state has the TTT and hys being at their lowest possible values of 0 seconds for TTT and 0 dB for hys. This starting state is expected to cause ping-pong's to occur because a handover will be triggered as soon as a neighbouring Base Station becomes better than the serving Base Station, instead of waiting to see if the neighbouring Base Station continues to be better for a period of time. The final starting start is to have the TTT start at a low value of 0.08 seconds and the hys start at a high value of 7.5 dB. This type of state was chosen to see how the well the system can optimise the values when it is likely that one will need to be increased and the other decreased.

5.2 Walking Speed

For the walking speed testing it is expected that the first scenario of having both the TTT and hys at their highest values will produce a large number of dropped calls, with the likelihood of ping-pong's occurring being very small, and prompt the optimisation system to reduce both values. For the second scenario with both the TTT and hys being middle values it is expected that the system will perform better than if no optimisation takes place. It is also expected that the system is likely to still reduced both values in an attempt to improve performance. For the third scenario with both the TTT and hys being their lowest possible values it is expected that no or very few dropped calls will occur and than ping-pong's are far more likely to happen. The system should compensate for this by increasing both values. In the final scenario where the TTT is a small value and the hys is a large one it is still expected that no ping-pong's will occur and that the ratio of called drops may be fairly high. It is expected that the system would attempt to improve the performance by increasing the TTT and decreasing the hys.

5.2.1 Large Starting Values

The results of how the optimised values compared to the static values can be seen in Figure 5.1 when both the TTT and hys started with their largest possible values of 5.12 seconds and 10 dB respectively. The results show that while the process of optimising the values initially generated a very large increase in the number of dropped calls. However, the system then managed to improve rapidly

and ended up having a better dropped call ratio by the end to the simulation run than that of the non-optimised system.




Figure 5.1: Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=5.12s hys=10dB when UE traveling at walking speeds.

Figure 5.2a shows how the TTT values were optimised over the course of the simulation run. It can be seen that all the base stations were in consensus that 5.12 seconds is too large of a values of the TTT and very quickly lowered the value. After this however there were effectively two different groups. The first group made up of base stations 1, 3, 4 and 5 ended up keeping the TTT value over 1 second. While the other group made up of base stations 0, 2, 6, 7 and 8 kept reducing the value to a most 0.48 seconds. It could be said that first group of base stations may not have optimised the TTT value as well as the second group because it can be seen that base stations 1, 4 and 5 kept switching between only two values when dropped calls occurred. It can also be seen the these changes happened often meaning that these base stations were still have dropped calls occurring, so it cannot be said that these values performed as well as those the base stations is the second group were using. While the changes made to the TTT value could be split up into two different groups the same cannot be said for the changes made to the hys value as seen in Figure 5.2b. It can be seen that while most of the base stations kept the hys value above 9 dB, base stations 0, 6, 7 and 8 lower the value more, with base station 8 settling between 5 and 5.5 dB. It can also be seen that base stations 1 and 3 appear to have gotten stuck between two non-optimal values as they keep switching between them and due to it happening often it means that dropped calls are still occurring so the system

should have switched to different values to see if they had performed better.




Figure 5.2: Illustration of how the TTT and hys values changed over time for large values when UE traveling at walking speeds.

From these results it can be said that the system performed as expected by reducing both the TTT and hys in the case of all the base stations. There were also a very high number for dropped calls and no ping-pong's, which was also expected in the simulation.

5.2.2 Middle Starting Values

Again as seen in Figure 5.3 the optimised values performed a lot better than the static values when they were originally set to the their middle values of 0.256 seconds for TTT and 5 dB for hys. The results also show that unlike with the results seen in Figure 5.1 the optimisation process did not greatly increase the number of dropped call at first. Instead it was the non-optimised values that began with the large number of dropped calls. This means that when these dropped calls began to appear the machine learning algorithm changed the value of TTT and hys in such a way that the dropped calls seen for the non-optimised values did not happen.




Figure 5.3: Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0.256s hys=5dB when UE traveling at walking speeds.

The way that the value of TTT trended during the simulation can be seen in Figure 5.4a. It can be seen that the majority of the base stations ended up having the TTT being 0.1 seconds with those base stations having very few changes from that value when they got there meaning that it could be an optimal value. As expected base station 4 came up with a different solution to the other base stations by actually increasing the value of TTT when the majority were reducing it.

Figure 5.4b shows that all but one base station were in consensus that the value of hys should be lowered from its starting value. Other than this though there was not very much consensus with the base stations finishing with any value in the range of 2.5 to 4 dB.

It is interesting to note that base station 8 actually ended up switching between the same states of values TTT=0.256 seconds hys=5 dB and TTT=0.16 seconds and hys=5.5 dB by the end of both simulations for starting at large and middle values. This means that these could be the most optimal states found for the base station.

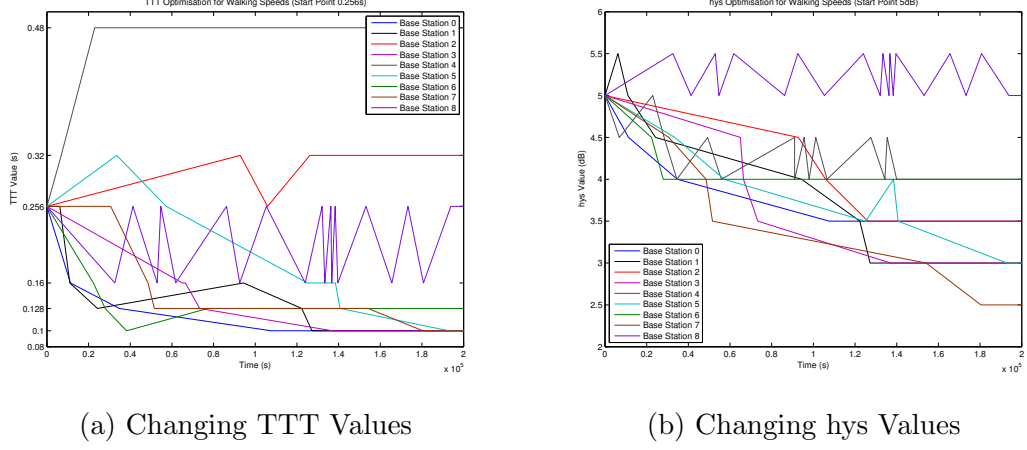


Figure 5.4: Illustration of how the TTT and hys values changed over time for medium values when UE traveling at walking speeds.

The system performed as expected in this scenario with the majority of the base stations reducing both the TTT and hys to attempt to improve the performance. Some base stations, however, took an unexpected approach by actually increasing the values.

5.2.3 Small Starting Values

It turned out that ping-pong's were a very rare occurrence in the simulation. This was most likely due to there being no fading in the simulation and that Random Direction mobility model would having the UE moving in one direction for a long time meaning that the only likely occurrence of a ping-pong would be if a handover took place and the UE then turned around and moved the other direction.

Figure 5.5 shows how the optimisation system performed against the static values when the simulation started with the TTT being 0 seconds and the hys being 0 dB. It can be seen that the optimised values only actually began to perform better than the static values at the very end of the simulation run. However, it can also be seen the ping-pong ratio for the static values was actually trending up while the ratio for the optimised values was trending down. This means that long term the optimised values should perform much better.

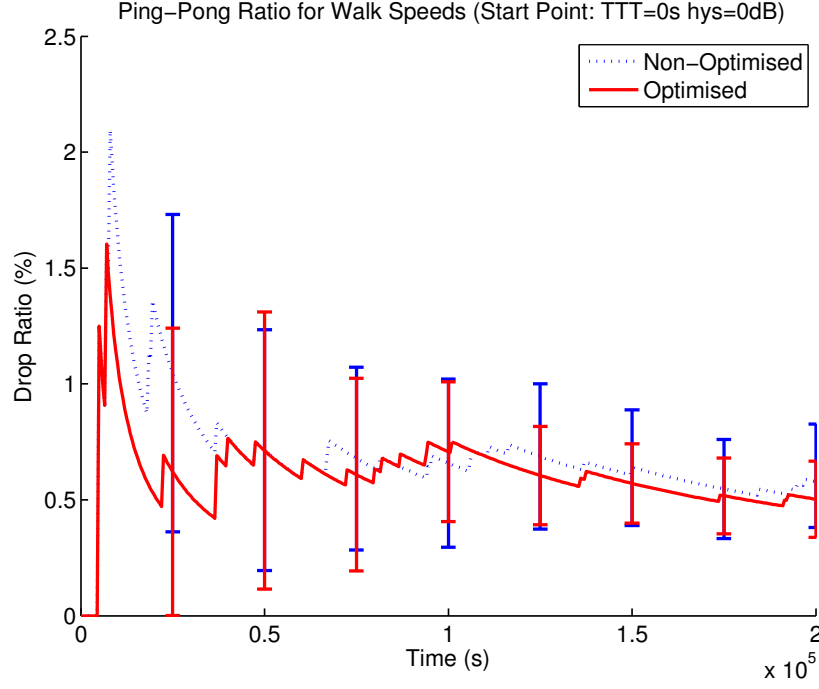


Figure 5.5: Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0s hys=0dB when UE traveling at walking speeds.

As for how the TTT and hys values changed over the simulation run in can be seen in Figure 5.6a that all the base stations were in consensus to keep to the TTT value at 0 seconds. This only changed at the very end of the run when base station 5 decided to increase the value to 0.04 seconds.

It can also be seen in Figure 5.6b that the majority of the base stations were in consensus that the hys should have been increase to 0.5 dB. However, this change was made at differing times because due to pong-pings happening so rarely and the optimisation system only making changes when something goes wrong.

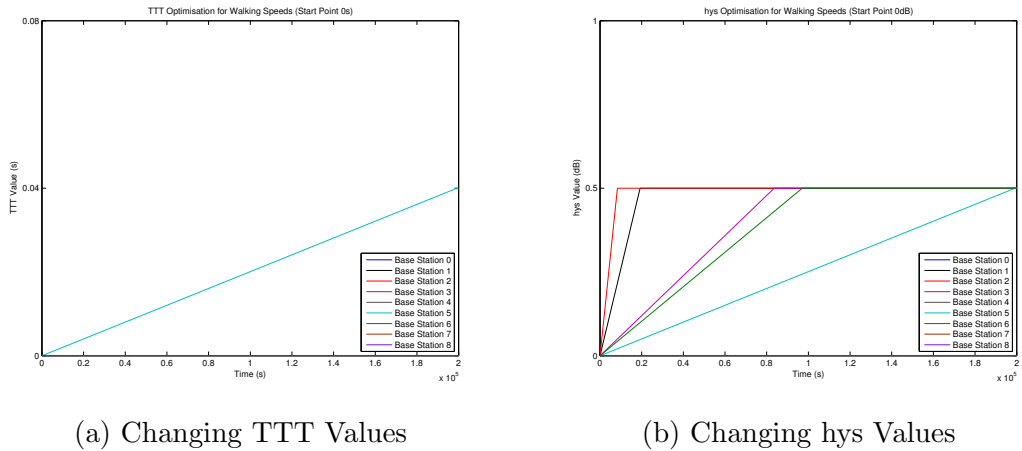


Figure 5.6: Illustration of how the TTT and hys values changed over time for medium values when UE traveling at walking speeds.

This scenario played out as expected with no dropped calls occurring; instead having ping-pongs determine the performance of the system. The system also compensated for this as expected by increase the value of the hys in most cases and by the end of the simulation run the performance of the optimised values had become better than that of the non-optimised values.

5.2.4 Large hys and Small TTT Starting Values

Much like the first three scenarios, it can be seen in Figure 5.7 that, the optimised values performed better than the static values when the TTT and hys started at 0.08 seconds and 7.5 dB respectively. Also much like in the second scenario the optimisation process did not begin by greatly increasing the number of dropped when trying to optimise the values.




Figure 5.7: Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0.08s hys=7.5dB when UE traveling at walking speeds.

The way in which the TTT values were changed during the optimisation process can be seen in Figure 5.8a. The figure shows that there was not very much consensus at all when it came to what value that TTT should be. It can also be seen that the value for the TTT for each base station was switching between the same two values a lot meaning that the optimisation system must have been making more changes to the value of the hys to try and get better performance.

Figure 5.8b shows that while the base stations changed their value of hys quickly after the simulation began there is no real consensus a lot for what value would be most optimal. The majority of the base stations ended up after their own unique value for hys by the end of the simulation run with the values being in the range of 3.5 to 7.5 dB. While there was no consensus on what the value for hys should be it can also be seen that when the base stations settled on their own value of hys they stayed with it meaning that there probably should have been more changes made to the TTT value to improve the performance.




Figure 5.8: Illustration of how the TTT and hys values changed over time for medium values when UE traveling at vehicle speeds.

In this scenario the system perform as expected in terms of decreasing the value of the hys. However, the system performed unexpectedly by also decreasing the value of the TTT when it was expected to actually increase it. Even with this the system still performed better than the simulation when run with static values.

5.3 Vehicle Speed

For the vehicle speed scenarios it is expected that number of dropped calls would increase over those seen in the walking speed scenarios due to the UE moving a lot faster and therefore giving a lot less time for the decision to handover or not. It is still expected that for the first scenario the system will reduce both values for TTT and hys. In the second scenario of the middle values it can also be expected the both values for TTT and hys will be reduced as dropped calls will still be expected fairly often at the speed the UE will be moving at. For the third scenario where both values for TTT and hys are at their lowest it is expected that the system will increase these values has ping-pong's should still be likely to

happen, even at the high speeds. In the fourth scenario it is expected that the system is more likely to reduce both values, unlike when the UE was moving at walking speeds, because of the speed the UE is moving at decisions need to be made quickly.

5.3.1 Large Starting Values

It can be seen in Figure 5.9 that the optimised system performed better than the system using the static values of 5.12 seconds for TTT and 10 dB for hys.




Figure 5.9: Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=5.12s hys=10dB when UE traveling at walking speeds.

Figure 5.10a shows how the TTT value for the base stations changed over time. It can be seen that the system did not appear to optimise the value very well. A lot of the base stations kept switching between the values of 5.12 seconds and 2.56 seconds, which for a UE moving at roughly 30 mph is far too high and this is shown by the fact that the base stations are switching back and forth between the values almost constantly. This same kind of results can also be seen for the base stations that lowered the values slightly more, they are still switching between values very regularly meaning that the performance could have been improved a lot more.

The same kind of results seen in the TTT values can also be seen in the hys in Figure 5.10b. The results show that while some base stations managed to reduce the value of the hys as far down as 6 dB, the majority of the base stations kept

the hys between 9 and 10 dB. This again looks to be far too high when the UE is moving at 30 mph as the base stations are constantly switching value but the performance never really improves.




Figure 5.10: Illustration of how the TTT and hys values changed over time for large values when UE traveling at vehicle speeds.

For this scenario while the system performed as expected and still performed better than that of the non-optimised system, it can be seen that it still did not perform all that well and appear be being getting stuck in non-optimal states.

5.3.2 Middle Starting Values

In this scenario the optimised system perform much better than of the static values, where it TTT and hys values were initialised to 0.256 seconds and 5 dB respectively. The graphs of the dropped calls ratios can be seen in Figure 5.11.






Figure 5.11: Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0.256s hys=5dB when UE traveling at vehicle speeds.

Figure 5.12a shows how the TTT value for each base station changed over the time the simulation was run. It shows that most of the base stations lowered the value for TTT while base station 6 greatly increased the value. Although this large increase appears strange the base station did not have a far larger number of changes when compared to the other base stations, which means that the values performed just as well as the other values chosen by the other base stations.

The way that the value for the hys was changed by the base stations can be seen in Figure 5.12b. It can be seen that the majority of the base stations kept the value of the hys around the starting point of 5 dB. Base station 4, however, reduced the value down to a minimum of 3.5 dB and be seen in have far less switches than the other base stations which means that it could have found a more optimal value.



(a) Changing TTT Values



(b) Changing hys Values

Figure 5.12: Illustration of how the TTT and hys values changed over time for medium values when UE traveling at vehicle speeds.

In this scenario it has been seen that the optimisation system functioned as expected with the majority of the base stations reducing both the TTT and hys values. However, there was a strange anomaly where one of the base stations greatly increased the TTT value, which was highly unexpected.

5.3.3 Small Starting Values

It can be seen that for this scenario the optimised system performed well and even though at the beginning of the simulation it performed worse, by the end the ratio of ping-pong's was for better than that of the non-optimised values.

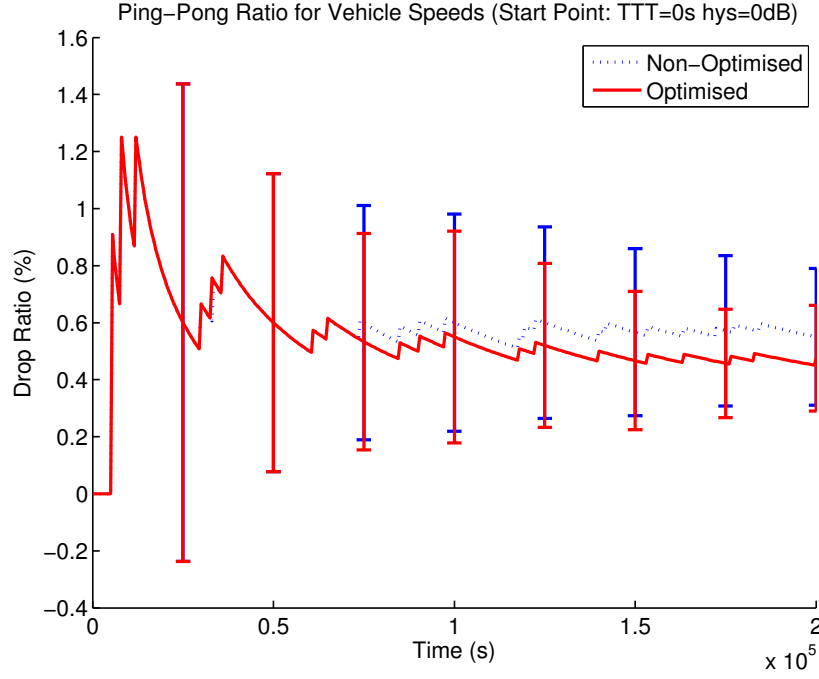


Figure 5.13: Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0s hys=0dB when UE traveling at vehicle speeds.

Figure 5.14a shows of the TTT value for base stations were changed by the optimisation system. It can be seen that most of the base stations stayed with the value being 0 seconds while other tried increasing the value to 0.04 seconds, with two of the base stations actually going back down to 0 seconds afterwards.

Figure 5.14b shows that the majority of the base stations increase the hys value 0.5 dB when a ping-pong occurred. Base station 3 actually increased the hys value to 1 dB when it had experienced two ping-pongs. Even with this it still looks like 0.5 dB may have been the most optimal value for the situation.




Figure 5.14: Illustration of how the TTT and hys values changed over time for medium values when UE traveling at vehicle speeds.

Again the optimisation system functioned as expected with increasing both the TTT and hys values due to ping-pongs occurring. This was expected because even though the UE is travel quickly the system still has to make sure that it does not trigger a handover too quickly.

5.3.4 Large hys and Small TTT Starting Values

As seen in the first three scenarios for the vehicle speed testing the optimised system managed to perform better than that of the non-optimised system when starting with the TTT being 0.08 seconds and the hys 7.5 dB. However, the improvement was not quite as great as seen in the other scenarios.

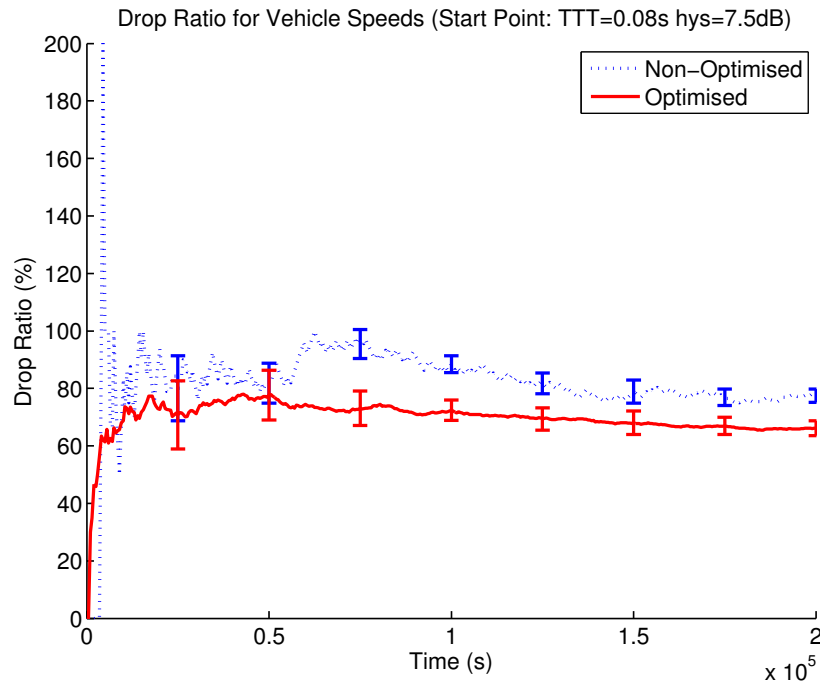


Figure 5.15: Graph of Optimised vs. Non-Optimised Results for Starting Point TTT=0.08s hys=7.5dB when UE traveling at vehicle speeds.

It can be seen in Figure 5.16a that the base stations did not really come to a consensus on what should be done to the TTT when a dropped call occurs. Some of the base stations increased the value while others lowered it to as low as 0 seconds. It is also hard to tell if it was correct to either increase or decrease the value as all the base stations are constantly switching between values which means calls are still being dropped.

Much like with the results seen for how the TTT was changed, there is no real consensus that can be seen for changed the value for the hys as seen in Figure 5.16b. It is seen that some base stations increased the value while other decreased it and only base station 0 decreased it by a large amount.

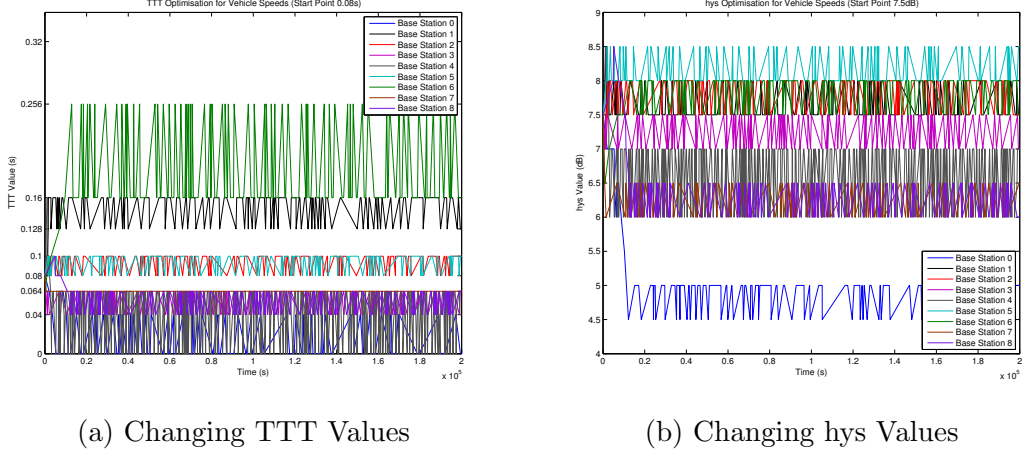


Figure 5.16: Illustration of how the TTT and hys values changed over time for medium values when UE traveling at vehicle speeds.

Much like with the other scenarios it is seen that while the system may have been an improvement over the static values the system still could have performed far better, as base stations appeared to keep switching between values that could be considered non-optimal.

5.4 Evaluation

Over the four scenarios used to test the optimisation system for walking speeds it can be seen that overall the system performed better than if no optimisation was done at all. It can also be seen that the system makes improvements to the performance very quickly. As seen in scenarios for the middle values and the large hys with small TTT the system managed to not have the large spike in dropped calls at the beginning of their simulation runs which were seen in their respective non-optimised runs.

In the scenarios for vehicle speeds it was seen that the system did performed better than if no optimisation was used. However, it was seen in the ways that the TTT and hys values were changed that they did not stop dropped calls constantly happening, which means that these values were likely non-optimal. Therefore it is possible that the system could have performed a lot better if the system did not get stuck between non-optimal states.

It can, however, be seen that when the system does not get stuck between non-optimal states that it will optimise the TTT and hys values as quickly as it is needed, i.e., whenever a dropped call or ping-pong occurs.

The optimisation system, however, also appears to have some drawbacks. It was seen in the first scenario that the optimisation system caused a very large increase to the dropped call ratio before improving it. This is a usual downside

in optimisation processes where 'things have to get worse before they can get better' and this process is a part of Q-Learning where the possible future rewards are taken into account when selecting a new state to move to. It was also seen in the scenarios using the vehicle speeds that the system appeared to keep getting stuck between two states that appear to be non-optimal which degrading the performance which had the potential to do a lot better.

Chapter 6

Future Work

Chapter 7

Conclusions

Bibliography

- [1] C. Cox, *An introduction to LTE: LTE, LTE-advanced, SAE and 4G mobile communications*. John Wiley & Sons, 2012.
- [2] 3GPP TS 23.002 V11.6.0, *Digital cellular telecommunications system (Phase 2+); Universal Mobile Telecommunications System (UMTS); LTE; Network architecture (Release 11)*, June 2013.
- [3] S. Feng and E. Seidel, “Self-organizing networks (son) in 3gpp long term evolution,” *Nomor Research GmbH, Munich, Germany*, 2008.
- [4] 3GPP TR 36.902 V9.3.1 Release 9, *LTE; Evolved Universal Terrestrial Radio Access Network (E-UTRAN); Self-configuring and self-optimizing network (SON) use cases and solutions (Release 9)*, May 2011.
- [5] T. Jansen, I. Balan, J. Turk, I. Moerman, and T. Kurner, “Handover parameter optimization in lte self-organizing networks,” in *Vehicular Technology Conference Fall (VTC 2010-Fall), 2010 IEEE 72nd*, pp. 1–5, IEEE, 2010.
- [6] 3GPP TS 36.331 V10.7.0, *LTE; Evolved Universal Terrestrial Radio Access (E-UTRA); Radio Resource Control (RRC); Protocol specification (Release 10)*, November 2012.
- [7] E. Alpaydin, *Introduction to Machine Learning*. MIT press, 2 ed., 2010.
- [8] A. G. Barto and R. S. Sutton, *Reinforcement learning: An introduction*. MIT press, 1998.
- [9] D. L. Poole and A. K. Mackworth, *Artificial Intelligence: Foundations of Computational Agents*. Cambridge University Press, 2010.
- [10] R. Roy, *Handbook of Mobile Ad Hoc Networks for Mobility Models*. Springer, 2010.
- [11] J. Chebil, A. K. Lwas, M. R. Islam, and A. Zyoud, “Comparison of empirical propagation path loss models for mobile communications in the suburban

area of kuala lumpur,” in *Mechatronics (ICOM), 2011 4th International Conference On*, pp. 1–5, IEEE, 2011.

- [12] N. Shabbir, M. T. Sadiq, H. Kashif, and R. Ullah, “Comparison of radio propagation models for long term evolution (lte) network.,” *International Journal of Next-Generation Networks*, vol. 3, no. 3, 2011.

Appendices

Appendix A

System Design and Evaluation

The contents...